Data Engineer

CV13 | INTERMEDIATE (SEMI-SENIOR)

CANDIDATE ID: CV13 | LEVEL: INTERMEDIATE (SEMI-SENIOR) | PROFILE: Data Engineer

PROFESSIONAL SUMMARY

Data Engineer with 5 years of experience designing, building, and maintaining scalable data pipelines and analytics infrastructure. Experienced with cloud data platforms (AWS, GCP), distributed processing (Spark, Airflow), and modern data stack tooling (dbt, Kafka). Security exposure is limited to data governance, encryption at rest/in transit, and GDPR/CCPA compliance requirements within data engineering contexts. No dedicated network security or cybersecurity background. Strong analytical mindset and engineering discipline.

EDUCATION

Bachelor of Science in Mathematics and Computer Science
Graduation Year: 2018

WORK EXPERIENCE

Senior Data Engineer
Digital Media & Advertising Company | Mar 2021 – Present (3.5 years)

- Designed and maintained data pipelines processing 5+ TB/day using Apache Spark on AWS EMR and Apache Airflow for orchestration.
- Built and maintained data warehouse architecture on Snowflake with dbt transformation layers serving 100+ analytics consumers.
- Implemented real-time streaming pipelines using Apache Kafka and Kafka Streams for user behavior event processing.
- Led data security review for CCPA compliance: implemented data retention policies, PII masking, and right-to-be-forgotten workflows.
- Managed AWS data infrastructure: S3, Glue, Redshift, Kinesis, Lambda.
- Collaborated with analytics and ML teams to design feature engineering pipelines for recommendation models.
- Mentored 2 junior data engineers and conducted code reviews.

Data Engineer
E-commerce Company | Jul 2018 – Feb 2021

- Built ETL pipelines using Python (Pandas, PySpark) and Apache Airflow.
- Managed BigQuery data warehouse on GCP.
- Created data quality monitoring using Great Expectations.
- Assisted with GDPR data mapping and anonymization requirements.

TECHNICAL SKILLS

Data Processing: Apache Spark (PySpark), Apache Kafka, Apache Airflow, Apache Flink (basic)
Data Warehousing: Snowflake (advanced), Google BigQuery, AWS Redshift, dbt (advanced)
Cloud Platforms: AWS (S3, EMR, Glue, Kinesis, Lambda, Redshift, Athena) GCP (BigQuery, Dataflow, Pub/Sub, GCS)
Languages: Python (advanced), SQL (advanced), Scala (basic), Bash
Data Governance: GDPR/CCPA compliance in data engineering, PII masking, data classification, data lineage (OpenLineage, Marquez)
Security (basic): Encryption at rest/in transit, IAM for data services, VPC design for data platforms (basic)
Visualization: Looker (basic), Tableau (basic)
Tools: Git, Docker, Terraform (basic), dbt Cloud

SOFT SKILLS

- Rigorous attention to data quality and pipeline reliability

- Strong collaboration with data analysts, scientists, and product teams
- Effective technical communication and documentation
- Problem-solving mindset for complex data architecture challenges
- Ability to manage competing priorities in a fast-paced environment
- Mentoring and knowledge sharing with junior team members

CERTIFICATIONS

- AWS Certified Data Analytics – Specialty — 2021
- Google Professional Data Engineer — 2020
- dbt Certified Developer — 2022
- Snowflake SnowPro Core Certification — 2021

LANGUAGES

- English: Full professional proficiency (C1)
- Spanish: Native
- Japanese: Basic (A1)

RELEVANT PROJECTS

Project 1: Real-Time User Behavior Pipeline (2022)
Designed and implemented a real-time event streaming pipeline processing 10M+ events/hour using Kafka and Kafka Streams, feeding real-time personalization models with sub-200ms latency. Reduced reporting lag from T+1 day to near real-time.

Project 2: CCPA Compliance Data Architecture (2022)
Designed and implemented CCPA-compliant data architecture covering data inventory, consent-aware pipeline suppression, PII tokenization, and automated right-to-deletion workflows across 12 data systems. Completed in time for regulatory deadline.

Project 3: Data Platform Migration to Snowflake (2023)
Led migration from legacy Redshift to Snowflake for a 300+ table data warehouse. Rebuilt 150 dbt models, implemented performance optimizations, and achieved 40% query performance improvement and 25% cost reduction.
